The motivation to vote: explaining electoral participation

1

The Decision to Vote or Not to Vote

Life is about making decisions, many small and a few big ones. Some decisions are made every day (what time should I get up?), others are made only once (or a few times) in a lifetime (what name should I give to my child?). Some decisions are made after a long process of reflection or deliberation (should I buy a house?) while others are made on the go, on the inspiration of the moment or gut feelings (should I have another beer?). Some are very personal (what should I eat at the cafeteria today?) and others are made jointly with others (what should we have for our Christmas dinner?).

The decision to vote or not to vote in an election can be deemed to be trivial. It is a small decision that we make rather infrequently and that has little or no consequence, for us as well as for society. The probability that a single vote will decide the outcome of an election is close to nil (Downs 1957; Mueller 2003; Owen and Grofman 1984), and so whether or not one votes will not decide which party will win the election.

Yet the decision to vote or abstain is not so inconsequential. First, it is not as infrequent as it may seem. In this century, that is, in the last nineteen years, the senior co-author of this book had to make that decision in seven (Canadian) federal elections, five (Quebec) provincial elections, and six (Montreal) local elections – once a year on average. This is in a country where there is no presidential or second chamber

election. In many places, there are also referenda, where the same decision to vote or not to vote needs to be made. The average citizen living in a democratic country where she acquires the right to vote at age eighteen and has a life expectancy of seventy-five years has the opportunity to vote in an election or referendum more than fifty times in her lifetime.

Whether we choose to vote or abstain tells us (and others) a lot about *who* we are. As we show later, the decision to vote or not to vote very much reflects what we like and do not like, in life and in society, and our values, particularly our conception of citizens' rights and duties in the polity. It is a decision that is affected not only by our self-identity, our feelings, and our beliefs but also by concrete cost-benefit instrumental considerations. It is both personal and social. It is affected by ethical views, even though our choices are not always consistent with our ethical aspirations. In short, it is a quintessential human decision, based on a combination of emotions and instrumental calculations, full of complexities, ambiguities, and sometimes contradictions.

Furthermore, the decision that most of us make most of the time, that is, to vote rather than to abstain, is paradoxical, in the sense that the rational person who calculates the personal benefits and costs of voting should come to the conclusion that she should abstain. She should abstain because the expected personal benefit of voting is extraordinarily tiny since the probability that her vote will decide the outcome is close to nil. Whether or not she votes will not decide who will be elected president or which party(ies) will form the government. Therefore, whenever there is some cost in voting, whether it is the time that it takes to go to the polling station and vote and/or the time to become informed in order to decide which party/candidate to support, the rational person should abstain (Downs 1957).

Yet turnout in national elections is typically around 70% (Blais 2018, using IDEA data), and so most citizens appear to make an "irrational" decision. This is known as the voting paradox. Our aim in this book is *not* to evaluate the merits and limits of rational choice theory (see Blais 2000). But the fact that most people appear to be "irrational," that

there is this apparent "paradox of voting," highlights the relevance of the question. There is no obvious answer to the question of why people vote. This is an enigma.

For all these reasons (and we suppose many others that we are unaware of), we decided to devote a good fraction of our time to addressing the turnout puzzle. We reveal how, after doing much research of our own, reading, discussing, and reflecting on the rich literature on the topic and related issues, we make sense of the simple act of voting or not voting. We present empirical evidence that supports our model. We argue that the decision to vote or abstain hinges on two basic predispositions (interest in politics and civic duty) and two election-specific judgments (caring and ease of voting). Clearly this is not an exhaustive model; many other attitudes come into play. Our claim is that with these four factors we can understand the basic motivations behind the turnout decision.

We focus on the individual-level determinants of turnout, that is, on the attitudes and judgments that lead someone to vote or abstain. There is a rich literature on the contextual factors that contribute to a higher or lower turnout (for a review and meta-analysis, see Blais 2006; Geys 2006; Cancela and Geys 2016; Stockemer 2017). Turnout, for instance, is higher when it is a close contest, when the office to be filled is more powerful, and when the previous election is not too recent (Blais and Dobrzynska 1998; Franklin 2004; Kostelka 2015). We do not deny the importance of these contextual factors (though many of the findings, especially with respect to the effect of the electoral system, do not seem very robust; see Blais 2006; Blais and Aarts 2006) but we wish to concentrate on the individual-level factors in this study.

It is possible that individual-level determinants of turnout vary across contexts (see Kittilson and Anderson 2011). We have no doubt that this is at least partly the case. For instance, the relationship between interest and politics is unlikely to be exactly the same in every country and in every type of election. Yet we start with the assumption that the motivations for voting and abstaining are basically the same in all elections, and that the impact of contextual factors is mostly

additive – that is, on top and independent of the individual-level factors. We revisit this assumption in Chapter 8.

The Framework

We construe the decision to vote or abstain as hinging upon the answers that each individual gives to four questions: (1) Do I like politics? (2) Do I have a duty to vote? (3) Do I care about the outcome? and (4) Do I find it easy to vote?

In the beginning, a person is either interested in politics or not. There are those who like politics and those who don't, just as there are some who do or do not like sports, arts, or religion. There are many reasons why people are more interested in some domains than in others. Our goal is not to understand why some people are interested in politics whereas others are not, though this would be a fascinating (and complex) study. Rather, we take this as a given, and look instead at how political interest, which we take to be equivalent to liking politics, affects the propensity to vote.

The basic intuition is simple and straightforward. If someone finds tennis exciting, she is prone to want to play and watch tennis and keep up with many events related to that sport. If someone finds it boring, then she sees little reason to follow it. The same rationale applies to politics – that is, some of us find it exciting and follow it passionately while others find it boring, complicated, or threatening, and they do not care about it or may even try to avoid it as much as possible (Hibbing and Theiss-Morse 2002).

These differences in taste are profound and enduring. Markus Prior (2010, 757), who has examined the stability of political interest over the life cycle with long-term panel data, arrives at the following conclusion: "Of the 58 stability coefficients ... for panel waves that occurred one year apart, only 10 have 95% intervals that do *not* include 1.0." Therefore, "people return to their stable long-term political interest levels quickly after perturbations caused by political or personal events. In short, political interest behaves like a central element of *political identity,* not like a frequently updated attitude" (763, emphasis added).

Being interested or not in politics is not a simple matter of taste; it defines who we are.

It is not farfetched to predict that those who are interested in politics are likely to vote and those who have no interest are inclined to abstain. There are of course other factors at play, but one's level of interest in politics, which remains remarkably stable over time, acts as a strong predisposition. Those who like politics like elections, and those who dislike politics dislike elections. The relationship is not perfect, as it is possible to be very interested in politics and to have little concern for a specific election, but we expect most of those who are interested in politics to find most elections exciting and to want to participate most of the time. The opposite should hold for those with little or no interest in politics.

This assumes that the driving force behind the decision to vote or abstain is *motivation* (hence the title of this book). The main reason, therefore, why many people do not vote is simply that they have little incentive to vote – they are not psychologically engaged. Contrary to Henry Brady and colleagues (1995), who argue that the main reason for lack of political participation is lack of resources, we assume that the main reason why some people do not participate in elections is that they do not want to – that is, motivation matters more than resources. The resource model is certainly relevant to the study of political participation broadly defined, but it is much less useful with respect to electoral participation, as Brady and colleagues (1995, 283) themselves acknowledge: "Indeed, political interest is much more important than resources if our main project is to explain voting turnout." Thus, the first question we should ask someone if we want to understand why she decided to vote or not is simply whether she likes politics.

The extent to which someone is interested in politics is not the only predisposition that matters. Quite a few people want to vote even though they are not interested in politics. The reason is their belief that they ought to vote, no matter how they feel about politics, elections, parties, or candidates. They believe that they have a moral obligation to vote, that is, they have a civic duty to participate in an election.

In a democracy, every citizen over a certain age has the right to vote. Whether citizens have a duty to vote is ambiguous. In countries where voting is compulsory, citizens have a legal duty to vote. Our study focuses on countries where voting is voluntary, but we should keep in mind that voting is officially defined as a legal duty to participate in many democracies (Singh 2019).

In those countries where voting is formally voluntary, the public discourse is ambiguous. While it is recognized that people have the right to abstain, there is the public norm that the *good* citizen has a civic duty to vote (provided it is not too complicated to do so; being sick, for instance, is a completely acceptable reason for abstaining). When asked in 1944 whether they see voting "more as a duty you owe to your country or more as a right to use if you want to," 59% of Americans chose "a duty" and 36% "a right" (Dennis 1970, 827). More recently, when asked how important it is for the good citizen to always vote in elections on a scale from 1 to 7, the mean score was 6.2, just slightly lower than the score for obeying the laws (Dalton 2008, 30). Moreover, about 90% of Canadians and 80% of British citizens agree with the statement that "it is the duty of every citizen to vote" (Blais 2000, 95).

The reasons why many people believe that they have a duty to vote are not always clear. One way to think about this is that people are motivated by reciprocity (Falk and Fischbacher 2006). The basic idea is that people wish to reward kind actions and punish bad ones. There is a huge experimental literature that supports the theory (see Dufwenberg and Kirsteiger 2004; Cox 2004). In the case at hand, when people are given a right that they cherish (the right to vote), they feel that they should reciprocate, and the most obvious way to reciprocate is to make use of that right. Another interpretation is that the civic norm of duty is learned at the community level. David Campbell (2006), especially, shows that the school and community civic climate at the time of adolescence affects adults' willingness to be politically active years later.

Whatever the reasons underlying the belief that there is a civic duty to vote, there is little doubt that many people subscribe to this view.

At the same time, we should not overstate support for this norm. It is politically correct to say that there is a civic duty to vote. We assume that indeed quite a few people strongly adhere to the view that they have a moral obligation to vote and that this strongly affects their decision to vote. But there are also many people who adhere to the opposite norm, that people are free to do what they want in a democracy and that there is nothing wrong with deciding not to vote. There are also many who do not have clear views either way, who pay lip service to the public norm when responding to a survey but who have not truly internalized the norm that citizens have a duty to vote in elections. The challenge is to sort out those who truly believe that they have a duty to vote.

It is useful to point out the similarities and differences in how political interest and sense of duty influence the decision to vote. Both are strong *predispositions* that people develop early in life and that are mostly stable over time. We have referred earlier to Prior's work (2010) demonstrating the remarkable stability of political interest. We do not have similar long-term panel data for sense of duty, but the evidence that we do have suggests strong stability. André Blais and Chris Achen (2019) report strong correlations over four waves, covering one year, during the 2008 American presidential election. Carol Galais and André Blais (2016a) find similarly strong over-time correlations in Spain over a period of eighteen months (four waves).

What is also common to political interest and sense of civic duty is that both are *general* predispositions that lead people to vote or abstain in any election. The other two considerations that we discuss below are more election-specific and vary over time, depending on the specificities of the context. Interest and duty act as broad attitudes that push individuals in one direction (voting or abstaining) in the absence of countervailing factors.

What distinguishes duty from political interest is that it is *moral*. The person who is interested in politics is inclined to vote because she *wants* to. The person who believes there is a civic duty to vote is inclined to go to the polling station because she feels she *ought* to. The interested person enjoys voting because politics is exciting. The dutiful citizen

decides to vote because her conscience tells her that she must do it; it is a duty that, like many other duties, needs to be fulfilled because it is the right thing to do. Duty therefore motivates one to vote in a different way than interest. The driving mechanism is normative rather than affective.

We expect those who are interested in politics to have a stronger sense of duty. After all, those who like politics have a positive prejudice about political matters, and they should be prone to think that people should participate in politics in general and particularly in elections. The relationship should be far from perfect, however. Some people are generally prone to think in terms of obligations, whereas others are deeply suspicious of so-called duties. People have different conceptions of what citizenship does and does not entail, and these conceptions are bound to shape their views about whether they have an ethical obligation to vote or whether it is a matter of personal choice.

The last two considerations that come into play in the turnout decision are at least partly election-specific. The first is how much the individual cares about the outcome. This "how much does it matter?" question corresponds to the *B* term in the rational calculus of voting model (Downs 1957; Riker and Ordeshook 1968). It boils down to whether, and how much, the person prefers one of the parties.[1]

As elegantly explicated by Anthony Downs (1957), in a two-party system, the citizen seeks to determine what each party will do over the course of the next mandate if it wins the election. If she believes that the two parties will adopt the same policies, she is indifferent and has no reason to vote. She is also indifferent if she expects the two parties to implement different policies but these policies are equally satisfactory or unsatisfactory, or if these differences concern issues that she does not care about. In short, the citizen votes only if she feels that the parties differ in meaningful ways about the issues that she is personally concerned with.

From this perspective, two conditions must be met for a person to care about the outcome of the election. First, she must care about the main issues that are debated in the campaign. Second, she must believe that the decisions that will be made about these issues depend to a

good extent on who will be elected. If either of these conditions is absent, the person is indifferent, and she has no reason to vote. It should be pointed out, however, that it is possible for a citizen to care a lot about the outcome of an election without paying much attention to the issues. This would be the case of voters who strongly prefer a given party, for example, because they trust its leader or simply because they identify with the party and are thus convinced that it is the best to govern the polity.

We expect this third consideration of our model, that is, "caring," to be positively correlated with political interest for three reasons. The first is that those who like and follow politics are more likely to be aware of differences between the parties. Those who do not follow politics regularly may be only vaguely aware of the positions of the parties and are unlikely to devote much effort to finding out these positions. The second reason is that those interested in politics are more prone to developing strong views and preferences, and thus to care a lot about what the government should and should not do. Again, the uninterested are more likely to care about who will win the football championship or who will win the Grammy Awards. The third reason is that the uninterested are more likely to distrust everything that politicians say during an election campaign. When you do not like politics, you are likely to dislike politicians and to be skeptical of their promises. The consequence is that the parties and candidates may all look alike.

The relationship between lack of general political interest and indifference in a specific election should be only moderately strong, however. Even the uninterested sometime get excited in a specific election and/or about a special issue, or they are attracted to a specific party or leader. Conversely, those who follow politics regularly may occasionally find little meaningful differences between the parties or they may not care about the issues discussed in a campaign. The relationship between duty and caring should be even weaker, especially once we take into account individuals' level of political interest. There should be some relationship since those who believe that they have a moral obligation to vote may also think that they have a moral obligation to cast an informed vote (see Blais, Galais, and Mayer 2019),

and so are more likely to search for at least some information about the parties. This should facilitate the process of forming an opinion about what the "good" and the "bad" options are. The dutiful person, however, does not have the inner motivation to follow politics closely (unless she is already very interested) and so is bound to pay little attention to an election campaign, and her preferences will often be relatively weak.

The fourth and last factor in our model is the perceived ease/difficulty of voting. This is the cost term ("C") in the rational calculus of voting. We start with the assumption that for most people most of the time voting is easy, and this explains in part why turnout is relatively high. It is precisely because voting is a simple, undemanding act that so many people believe they have a moral duty to vote. Most people would find it unreasonable to require all citizens to participate actively in groups, organizations, parties, or demonstrations, but should not the "good" citizen do her part and contribute a little bit of her time and go to the polling station?

Going to the polling station is easy for most people, but it is difficult or complicated for some.[2] It is not easy for all those who happen to be sick or away from home on the day of the election. It may not be easy for all those who suffer from some handicap. For some it may be a source of stress. In that case, going to the polling station as such is easy, but they may not know or understand exactly what they are supposed to do, what the ballot paper looks like, and what they are supposed to write on that ballot. Even in our advanced societies there are many people who are practically illiterate; for many of them, the act of voting is far from simple. In those circumstances, staying home has a lot of appeal.

Since Downs (1957), researchers distinguish between the cost of going to the polling station and the cost of looking for information to help one make up one's mind which party to vote for. This distinction is not as straightforward as it may first appear. The cost of finding information about the parties is clearly higher for those who do not like politics and do not follow the news, but the crux of the problem is that such people lack the motivation to follow politics to begin with;

they are not interested in politics. Similarly, those who have formed strong preferences about the issues and the parties face little information cost, but this is simply because they care a lot about the outcome of the election.

In this study, we do not distinguish between voting and information costs (but see Blais et al. 2019). We simply rely on people's overall subjective judgment about how easy or complicated they find voting to be. Because it is a subjective perception, it is likely to be shaped in part by people's predispositions, most strongly their interest in politics and secondarily their sense of civic duty, as well as by how much they care about the outcome. Still, these correlations should not be very strong, since most of us sometimes find ourselves in situations where we are sick, depressed, or overwhelmed with more pressing concerns. And it is fair to predict that whenever the cost of voting becomes high, the temptation to stay home becomes strong.

Our main goal is to show that these four basic factors (interest, duty, caring, and ease) help us a lot in making sense of the decision to vote or abstain. The first three are clearly motivational factors and are thus perfectly in line with our motivational account of the turnout decision. Ease of voting needs to be incorporated as an additional factor, but it is clearly as individual and subjective as the first three factors.

We then devote an entire chapter to an alternative explanation of turnout that has gained popularity in the recent literature, namely, that voting is a habit (see especially Franklin 2004). This interpretation is based on the accurate observation that the decisions that people make to vote or abstain in different elections are not independent. A person who voted (or abstained) in the previous election is likely to vote (or abstain) in the next one. From this observation, it is a small step to inferring that voting is a habit that people display repeatedly over time.

We do not dismiss the possibility that voting is a habit for some people, but it is strange to have a habit involving something that occurs relatively infrequently and irregularly. We pointed out earlier that people typically decide whether or not to vote about once a year. This

is not frequent enough to develop a habit, we believe, especially since the timing is irregular, and the rules and the set of options (the parties) often vary across types of elections.

Furthermore, and most importantly, people may well repeat the same behaviour over time simply because the values and/or attitudes that drive their behaviour remain the same. Someone who is very interested in politics or has a strong sense of duty is prone to vote in every election, especially as interest and duty are unlikely to vary substantially over time. The opposite holds if the person is uninterested in politics and does not believe that she has a moral obligation to vote. As Eric Plutzer (2002) notes, it is important to distinguish between persistence and inertia (habit). We should point to the presence of habit only if we can show that the turnout decision in an election depends on the turnout decision made in the previous election, not on the factors that shaped the initial decision. We perform more appropriate tests of the habit hypothesis, and find little support for it.

Finally, we consider contextual effects. Traditionally, there have been two separate streams of research on turnout (Blais 2006). One stream is based on aggregate turnout data and focuses on the contextual-level factors (mostly institutional) that are associated with lower or higher participation rates. The second is based on survey data and examines the individual-level factors (mostly attitudinal) that are associated with the propensity to vote. More recently, an important new stream of research has attempted to combine these two approaches. With the advent of large datasets combining survey data from many different elections and countries, through such initiatives such as the Comparative Study of Electoral Systems (CSES), it has become possible to estimate within the same model the effect of both individual and contextual factors. In this context, special emphasis has been placed on discovering *interaction* effects, whereby the impact of individual variables is *conditional* on contextual factors (see Anderson and Dalton 2011).

We see a lot of merit in this new approach. The Making Electoral Democracy Work (MEDW) data that we use in this study are precisely based on the idea that it is crucial to look at how the rules of the game (the institutions) affect not only parties' and voters' behaviour but also

how they make up their minds – that is, the considerations that shape their decisions.

Yet we claim that the considerations that lead people to vote or not to vote are basically the same in all elections. We therefore start with a simple model that is tested with a merged dataset that includes all the elections covered by MEDW data. In the last chapter, we explicitly test for interaction effects between the four individual-level attitudes and contextual variables. We do find some context-specific patterns, but we show that they are relatively rare and that their impact is, in substantial terms, modest. We conclude that it is fair to assert that in established democracies people decide to vote or not to vote for similar reasons across all kinds of contexts.

The Approach

For most of the analyses, we use the MEDW surveys that were conducted in five countries between 2011 and 2015 (Blais 2010a; Stephenson et al. 2017): Canada, France, Spain, Switzerland, and Germany. These countries were chosen mainly because, although all are by now established and developed democracies, they differ most especially with respect to electoral system. Canada and France both have "majoritarian" rules for their national elections, Canada a first-past-the-post system and France a two-round system; Spain and Switzerland have proportional representation, while Germany has a compensatory mixed voting system.

We do not claim that these countries constitute a representative sample of established democracies. Europe is overrepresented, but clearly Europe dominates the list of established democracies, especially those with voluntary voting. There is an overrepresentation of federal and relatively decentralized countries, as well as of countries with supranational elections (mostly due to the overrepresentation of Europe). These biases should be kept in mind.

What is more important is that these countries represent a great variety of contexts. Some countries (France and Germany) are quite large in terms of population, and one (Switzerland) is very small. At the time of our study, two countries (France and Spain) were in deep

recession while the economic situation in the other three countries was relatively good. As mentioned above, these countries have very different voting systems. By combining these diversified cases, we hope to highlight the common patterns that emerge when it comes to deciding whether to vote or not. At the same time, it should be clear that our interpretation is confined to established democracies where voting is not compulsory. How the model would need to be modified in the case of non-established democracies (or non-democratic elections) and/or when voting is made compulsory is explored in the conclusion.

Turnout in the five countries considered here is somewhat lower than the average in contemporary democracies for national lower house elections, which is about 70% of registered electors (Blais 2018). Turnout in the MEDW national lower house elections was 49% in Switzerland (2011), 55% in France (2012), 69% in Canada (2015) and Spain (2011), and 72% in Germany (2013). The median turnout for the national lower house election covered by the study was 69% (very typical), but the mean is 63% (somewhat low).

In each of these five countries, we selected two regions: Quebec and Ontario in Canada, Lower Saxony and Bavaria in Germany, Zurich and Lucerne in Switzerland, Catalonia and Madrid in Spain, and Île-de-France and Provence (sometimes labelled "PACA" for Provence à Côte d'Azur) in France.[3] We selected regions that differ in their party systems, with the constraint that the region had to be populous enough that we could obtain a relatively large sample (about 1,000) of respondents in each case. We selected two regions within each country because we wanted to compare subnational and national elections.

We thus have ten cases, that is, two regions in each of our five countries. In two countries (France, Germany, and Spain), we cover three separate elections: supranational (the European 2014 election), the national election for the lower house, and subnational.[4] In Canada, Spain, and Switzerland, we examine the national and the most important subnational election, which we call "regional." In Canada, these

are provincial elections, in Switzerland cantonal elections, in Germany state (Lander) elections, and in Spain (autonomous) regional elections. In the case of France, we selected municipal instead of regional elections because the former are generally considered to be more important, as indicated by their higher turnout rate.[5]

In the two French regions and in Lower Saxony, we have three elections, while in the two Canadian provinces, the two Swiss cantons, the two Spanish regions, and the Bavarian state there are two elections. And we have the additional case of British Columbia for the 2015 Canadian election. This yields a total of twenty-four elections in eleven different regions. Note that in the case of national and supranational elections, the two elections that we examine in two different regions are part of the same election. We therefore have seventeen "independent" elections. As the region is the common unit of analysis, we systematically refer to twenty-four elections.[6]

Except for Bavaria and the 2015 Canadian election, the MEDW survey consisted of a pre-election wave with about 1,000 respondents (usually in the last ten days of the campaign) and about 750 respondents (out of the initial 1,000) in the post-election wave (usually in the seven days following the election). The pre-election wave took about twenty minutes and the post-election wave about ten minutes. We use the post-election wave for the dependent variable (whether the person voted or not) and the pre-election wave for the main independent variables (the four attitudes).

In the case of the 2015 Canadian election, we drew larger samples in each of the three provinces. In the end, we had 1,879, 1,891, and 1,849 respondents in the pre-election wave in British Columbia, Ontario, and Quebec, respectively, and 1,195, 1,308, and 1,206 in the post-election wave. In the case of Bavaria, we have a special three-wave panel, with the first wave occurring just before the September 15, 2013, regional election, the second wave right after and right before the September 22 national election, and the third and last wave immediately after the national election. The sample sizes for these waves were 4,261, 3,575, and 2,895, respectively.

Table A.1 of Appendix 2 summarizes the information about each of the twenty-four elections. The mean turnout in the twenty-four elections is 61% and the median 58%. The five Canadian elections are single-member district plurality elections, while the Swiss and the Spanish elections are proportional representation (PR) elections. The German national and state elections are mixed compensatory (with two votes) while the German European elections are PR. Finally, the voting system varies across the three levels in France: the French European election is held under PR, the national legislative election is single-member district two-round (majority/plurality), and municipal elections are held under a two-round mixed system that guarantees the winning list an absolute majority of seats, which can be considered mixed majoritarian.[7]

The main dependent variable throughout the book is whether the person voted or not. This information is provided in the post-election survey. In fact, there was a wording experiment. The first sentence of the turnout question was identical for all respondents: "In each election we find that a lot of people were not able to vote because they were not registered, they were sick, or they did not have time." This sentence is meant to make it easier for people to admit that they did not vote.[8]

Half of the respondents (the control group) were then simply asked, "Were you personally able to vote in this election?" with the response categories being "yes," "no," and "don't know." The other half (the treatment group) were asked instead, "Which of the following best describes you?" with the response categories being: "I did not vote in the election," "I thought about voting but didn't this time," "I voted in the election," and "don't know." In both cases, we assume that "don't know" corresponds with abstaining. As expected, the treatment version facilitates the admission of abstention and yields a lower turnout (Morin-Chassé et al. 2017). We have merged the two versions in all the analyses reported below. In Table A.2 of Appendix 2, we show that the patterns remain the same with the two versions of the question. The only interaction that is significant is related to "care." The substantial difference should not be overstated, however. Concretely, going from 1 standard deviation below the mean in care to 1 above the mean

has an impact of 13 percentage points in the treatment group compared with 8 percentage points in the control group, for a net difference of 5 points.

The MEDW data are based on online quota-based surveys that guarantee the representativeness of the samples with respect to gender, age, education, and region. Like almost every survey, the reported turnout is much higher than the official turnout. This is so first and foremost because of a self-selection bias. Those who are more interested in politics (and more inclined to vote) are more prone to agree to participate in a study that deals with politics. There is, on top of that, a social desirability effect. There is the public norm that the good citizen should feel a moral obligation to vote. For this reason, some people are reluctant to admit that they did not vote, and thus some abstainers indicate that they voted. The consequence is that abstention is almost always underestimated in surveys, and the MEDW surveys are no exception to this rule.

Most of our analyses, as is usually the case, are based on self-reported vote. We would of course prefer to have validated vote, as there is an overrepresentation of respondents saying that they voted (Rogers and Aida 2014; Selb and Munzert 2013). Does this introduce a major bias? We do not believe so. In a recent study, Chris Achen and André Blais (2016) use the 1980, 1984, and 1986 American National Election Studies (ANES) to examine the correlates of intention to vote, reported vote, and validated vote. They look at the impact of age, education, interest, duty, care, and party identification on intended, reported, and validated vote. They find that "all of the substantively relevant variables are statistically significant and with the correct sign in all three equations" (200), and that "reported vote is better than intended vote as a proxy for actual turnout" (205); they conclude that "our findings support a circumscribed and qualified endorsement of the current practice of conflating studies of intended vote, reported vote, and validated vote" (207). We recognize, however, that the effects reported in this study are likely to be slightly overestimated.

Throughout the book, we use turnout-weighted data; that is, the data are weighted so that the reported turnout in the survey corresponds

FIGURE 1.1 The funnel of causality

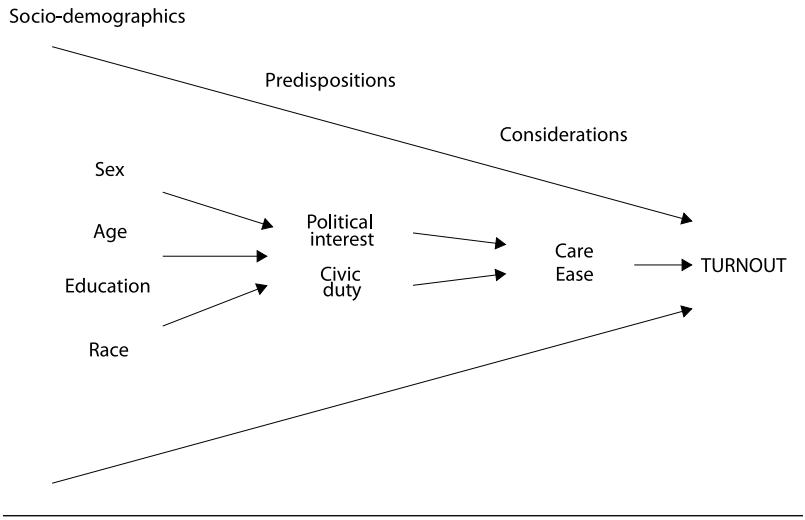

to the official turnout. We do this mostly for descriptive purposes. When we present frequencies or predicted probabilities of voting across different groups, the turnout figures are more realistic this way. This rests on the assumption that the differences in our sample between voters and abstainers are similar to the differences that exist in reality between these two groups.[9] We are assuming that this assumption is approximately correct. In Table A.3 of Appendix 2 we show that the patterns are very much the same with unweighted data.[10]

The analysis is based on a funnel of causality approach (see Campbell et al. 1960; Miller and Shanks 1996; Blais et al. 2002). We start with the most distant factors and then move to consider the more proximate causes. We first examine the socio-demographic correlates of voting, then the two basic predispositions (interest and duty), and finally the two election-specific considerations (care and ease). These can be seen as three distinct "blocs." (See Figure 1.1.)

In each case, we first present descriptive information about the factor (for example, duty) that is the focus of the chapter. Second, we look at the correlations between this factor and the antecedent

variables (in the case of duty, socio-demographic characteristics and political interest). The factor is then analyzed as the dependent variable. In a third step, the goal is to ascertain how that factor, now construed as independent variable, affects the propensity to vote, controlling for the antecedent variables. In all these analyses, all the variables are at the individual level, but we are also controlling for the specificities of each case using elections fixed-effects as we include twenty-three dummies for each election (except the regional election in Lower Saxony, which is the reference category).[11]

The approach is different in the last two chapters of this book. In Chapter 7, we take up an alternative interpretation that voting is in good part a habit. We review the various studies that have attempted to test the habit hypothesis and we show that these studies are not very satisfactory. We propose another test of that hypothesis, based on the plausible assumption that the propensity to have a habit is strongly correlated with age. We infer that if people vote or abstain out of habit, the determinants of the turnout decision should vary over the life cycle; that is, the turnout decision should be more strongly affected by values and attitudes among the youth, who presumably have not yet acquired a habit. Using three different datasets, we show that this is not the case, and we conclude that the habit hypothesis is not compelling. Moreover, we also use a different proxy for habit – whether the respondent always voted or abstained and we still find no evidence to support the habit interpretation.

Chapter 8 deals with contextual effects. As indicated above, the focus in this study is on individual-level determinants of the decision to vote or abstain. Clearly, however, this decision is also affected by contextual-level factors. Our claim is simply that these contextual variables correspond to additional causes of turnout that should be considered in a complete account, but that our basic individual-level model that is explicated in this book accounts for most of the variation. Indeed, as we show in Chapter 8 using a multi-level (mixed-effects logistic regression) model, 94% of the variance is due to differences across individuals and only 6% to contextual differences across the twenty-four elections.[12]

We also pay attention to potential interaction effects, that is, whether interest, duty, care, and ease matter more or less in specific contexts. We do find some interesting interaction effects but they are the exception rather than the norm, and, most importantly, they are very modest. This finding buttresses the claim that our model about the factors that drive the decision to vote or abstain applies in all contexts, at least in well-established democracies where voting is not compulsory.

Most of the empirical evidence that we present is based on the MEDW data. The reason is obvious. As this research was directed by the senior co-author, the survey questionnaires included questions designed to tap each of the four major variables that our motivational model incorporates. Furthermore, we believe that it is crucial to test our model with data collected in many different countries and in many different types of elections. The MEDW data satisfy these two criteria. Using the same dataset throughout the book makes it easier for readers to see how the various findings mesh together.

We do, however, use other datasets when they are required to provide more robust tests of our model. This is particularly the case in the last chapters, when we examine the role of habit and contextual factors that are outside our model.

The analyses presented here are based on cross-sectional survey data. Because our model focuses on the motivation (or lack thereof) to vote, we need to tap citizens' attitudes and relate them to their turnout decision, and survey data are therefore essential. As a consequence, experimental studies are not appropriate for testing our model unless they are complemented by survey data, which is seldom the case.[13] We discuss the limitations of experimental research in this specific respect in Chapters 4 and 7 on duty and habit. That being said, it would be better to have longitudinal panel survey data than cross-sectional survey data. Unfortunately, longitudinal panel survey data about the motivation to vote are almost non-existent, so we have to do with the MEDW cross-sectional survey data.[14]

The two main risks associated with the use of cross-sectional data are the possibility that the observed relationships may be spurious and

the possibility of reverse causation. Our model assumes that the decision to vote is driven by two strong predispositions, political interest and sense of civic duty. The risk that the observed relationships between interest or duty and turnout are spurious is reduced if these two attitudes are formed early in life and do not change much over the life cycle. In Chapters 3 and 4, we present and discuss evidence provided by longitudinal panel surveys that support our claim that these two attitudes are indeed quite stable. We also refer to studies that show little rationalization from turnout to duty and interest, that is, there is little evidence that voting makes people more interested in politics.

Finally, with respect to care and ease of voting, the risk of spuriousness is small since we control for powerful predispositions (political interest and duty) as well as age and education. We cannot rule out the possibility of rationalization, though that risk is reduced by the fact that these considerations are measured in the campaign survey.

We would of course prefer to test our model with longitudinal panel survey data, and we hope that future research will move in that direction. We believe, however, that the limitations of the cross-sectional data that we utilized are mitigated in this case because there is good empirical evidence that the attitudes that are at the beginning of the funnel of causality constitute strong and stable predispositions.

Our goal is therefore to propose an elegant and parsimonious model of the individual decision to vote in an election, to show that the MEDW data support that model, and to demonstrate that prior research provides additional support for our argument. We begin our empirical investigation by answering a simple question: Who votes?

2

Who Votes?

As indicated in Chapter 1, our goal is to explain why people do or do not vote. It is not possible, however, to provide a sound explanation of why people do what they do if we do not get the facts right. A good description is a necessary (though not sufficient) condition for a good explanation. We thus begin with some descriptive patterns before we consider the causes of electoral participation.

Indeed, the first question that comes to mind when we study turnout at the individual level is: *Who* votes? "Who" refers to the socio-demographic characteristics of voters and abstainers. Is turnout associated with age, race, gender, education, or income? We wish to know the socio-economic profile of voters and abstainers for many reasons. The first is simple curiosity. We want to know what kinds of citizens are more likely to vote or abstain, just as we are curious to learn what kinds of people get up early or late, are vegetarians, or listen to jazz. Second, knowing who votes should help us understand why they vote or should make us become skeptical about some interpretations. For instance, the fact that the better educated are more likely to vote (see below) makes us suspicious about the rational choice model of turnout, since the better educated should better understand that their own single vote is extremely unlikely to make any difference. The third reason is normative. If some groups are much less prone to vote than others, then voters are a biased sample of the eligible population, and we may legitimately wonder about the biases that this introduces

in the policy-making process (see Leighley and Nagler 2014, chap. 6). The fourth reason is that when we examine the impact of attitudes on the decision to vote, as we do in subsequent chapters, we need to control for possible spurious effects, and the most obvious antecedent causes of these attitudes are precisely socio-demographic characteristics.

We focus on two socio-demographic characteristics: age and education. We do so because prior research has shown that these are the two strongest socio-economic correlates of turnout. In their classic book *Who Votes?* Raymond Wolfinger and Steven Rosenstone (1980, 102) conclude that the two most important correlates of voting in the United States are education and age. The more recent *Who Votes Now?* (Leighley and Nagler 2014) pays more attention to income, but the authors recognize that "education still trumps income as a predictor of turnout" (66). They also indicate that the "age-related patterns in turnout confirm previous findings," though they note increases in turnout in the 2004 and 2008 elections among both the youngest and oldest age groups.

These studies deal with the United States, which is in many ways a special case. André Blais (2000) examines the socio-demographic correlates of voting in a merged dataset covering nine countries. He looks at the impact of age, gender, education, income, religiosity, marital status, union membership, and employment (being unemployed, retired, or housewife), and finds that "the two most crucial socio-economic determinants of voting are education and age" (52). All other variables have much weaker effects. Furthermore, Neil Nevitte and colleagues (2009) perform a similar analysis, covering twenty-three countries and thirty-two elections. They include age, education, income, marital status, church attendance, place of residence, unionization, gender, and employment status. The two variables that are statistically significant in the greatest number of elections are age and education.[1]

This does not mean that age and education are the most powerful predictors of turnout in each election. Indeed, Aina Gallego (2015) shows that the educational gap in turnout varies substantially across countries. Our more modest claim is that age and education are the

two socio-economic variables with the strongest and most consistent relationships with turnout.

Age

Concentrating on these two socio-demographic characteristics in our analysis enables us to examine their relationship with turnout in greater detail. Wolfinger and Rosenstone (1980) show that the relationship is non-linear, that is, turnout starts declining at over 70 years of age. They point out, however, that once education, sex, and marital status are controlled for, there is no real decline at old age. Furthermore, Leighley and Nagler (2014) find that in more recent elections in the United States the turnout rate of the 76–84 age group is in fact higher than average. Governmental (register-based) data in Denmark indicate, however, that turnout peaks at around 65 years, and then starts declining (Bhatti and Hansen 2012). According to Elections Canada (2012) estimates, the drop in Canada begins only at around age 75. Furthermore, there is some evidence, again based on register-based data, that turnout declines between the ages of 18 and 21, as many people leave their parents' home, and so the relationship between age and turnout is a sort of roller coaster (Bhatti, Hansen, and Wass 2012; Bhatti and Hansen 2012).

To complicate things further, it is not always clear how to interpret age differences in the propensity to vote. The usual interpretation, emphasized by Wolfinger and Rosenstone (1980), is the life-cycle one. This makes a lot of sense. As people grow older, they become more integrated into social and political life, and they are more prone to vote (Milbrath and Goel 1977). At the other end of the life cycle, they may become less integrated (especially if they become widowed), and their level of participation declines.[2]

People of different ages belong to different generations. There is strong evidence that the turnout decline that has occurred in most democracies around the end of the twentieth century is in good part a generational effect, that is, the turnout rate of youngest citizens is now considerably lower than it used to be. Blais and colleagues (2004)